

AI19542 - DATA SCIENCE USING R - LAB MANUAL



RAJALAKSHMI
ENGINEERING COLLEGE
An AUTONOMOUS Institution
Affiliated to ANNA UNIVERSITY, Chennai

**DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND MACHINE
LEARNING**

AI19542 – DATA SCIENCE USING R

LAB MANUAL

THIRD YEAR

FIFTH SEMESTER

2024 - 2025

ODD SEMESTER

Ex No:1

Date:

Basics of R – data types, vectors, factors, list and data frames

AIM:

To implement and understand the basics of R programming with its data types, vectors, factors, list and data frames.

ALGORITHM:

1. Start
2. Assign values in logical, numerical, character, complex and character in raw form to a variable v.
3. Print the class of v.
4. Assign a vector for subject Names, temperature and flu_status for three patients using c() function and access the elements.
5. Create a factor using factor() with duplicate values and assign level with distinct values.
6. Display the specific element and check for certain values in factor.
7. Create a list using list() from the patient details and access the multiple elements.
8. Create a data frame using data.frame() with multiple vectors as features. Access the elements.
9. Create a matrix using matrix() with different allocations and access the elements.
10. Stop.

PROGRAM:

#Data Types

v<-TRUE

print(class(v))

v<-23.5

print(class(v))

v<-2L

print(class(v))

v<-2+5i

print(class(v))

v<- "TRUE"

print(class(v))

v<-charToRaw("Hello")

print(class(v))

#Vectors

subject_name<-c("John Doe","Jane Doe","Steven Grant")

temperature<-c(98.1,98.6,101.4)

flu_status<-c(FALSE,FALSE,TRUE)

temperature[2]

temperature[2:3]

temperature[-2]

#Factors

gender<-factor(c("MALE","FEMALE","MALE"))

gender

blood<-factor(c("O","AB","A"),levels=c("A","B","AB","O"))

```

blood[1:2]
symptoms<-factor(c("SEVERE","MILD","MODERATE"),
                 levels=c("MILD","MODERATE","SEVERE"),
                 ordered=TRUE)
symptoms>"MODERATE"

#Lists
subject1<-list(fullname=subject_name[1],
               temperature=temperature[1],
               flu_status=flu_status[1],
               gender=gender[1],
               blood=blood[1],
               symptoms=symptoms[1])
subject1
subject1[2]
subject1[[2]]
subject1$temperature
subject1[c("temperature","flu_status")]

#Data Frames
pt_data<-data.frame(subject_name, temperature, flu_status,
                    gender,blood,symptoms)
pt_data
pt_data$subject_name
pt_data[c("temperature","flu_status")]
pt_data[c(1,2),c(2,4)]
pt_data[,1]
pt_data[,]

#Matrices
m<-matrix(c(1,2,3,4),ncol=2)
print(m)
m<-matrix(c(1,2,3,4,5,6),nrow=3)
print(m)
print(m[1,])
print(m[1,])
thismatrix <- matrix(c("apple", "banana", "cherry","orange"), nrow = 2, ncol = 2)
for (rows in 1:nrow(thismatrix)) {
  for (columns in 1:ncol(thismatrix)) {
    print(thismatrix[rows, columns])
  }
}

```

OUTPUT:

```
File Edit Selection View Go Run Terminal Help
PROBLEMS 73 OUTPUT DEBUG CONSOLE TERMINAL JUPYTER

[1] "logical"
[1] "numeric"
[1] "integer"
[1] "complex"
[1] "character"
[1] "raw"
[1] 98.6
[1] 98.6 101.4
[1] 98.1 101.4
[1] MALE FEMALE MALE
Levels: FEMALE MALE
[1] O AB
Levels: A B AB O
[1] TRUE FALSE FALSE
$fullname
[1] "John Doe"

$temperature
[1] 98.1

$flu_status
[1] FALSE

$gender
[1] MALE
Levels: FEMALE MALE

$blood
[1] O
Levels: A B AB O

$symptoms
[1] SEVERE
Levels: MILD < MODERATE < SEVERE

$temperature
[1] 98.1

[1] 98.1
[1] 98.1
$temperature
[1] 98.1

$flu_status
[1] FALSE

subject_name temperature flu_status gender blood symptoms
1 John Doe 98.1 FALSE MALE O SEVERE
2 Jane Doe 98.6 FALSE FEMALE AB MILD
3 Steven Grant 101.4 TRUE MALE A MODERATE
[1] "John Doe" "Jane Doe" "Steven Grant"

temperature flu_status
1 98.1 FALSE
2 98.6 FALSE
3 101.4 TRUE
```

```
File Edit Selection View Go Run Terminal Help BasicsO
PROBLEMS 73 OUTPUT DEBUG CONSOLE TERMINAL JUPYTER

2 98.6 FALSE
3 101.4 TRUE
temperature gender
1 98.1 MALE
2 98.6 FEMALE
[1] "John Doe" "Jane Doe" "Steven Grant"
subject_name temperature flu_status gender blood symptoms
1 John Doe 98.1 FALSE MALE O SEVERE
2 Jane Doe 98.6 FALSE FEMALE AB MILD
3 Steven Grant 101.4 TRUE MALE A MODERATE
[1] [,1] [,2]
[1,] 1 3
[2,] 2 4
[1] [,1] [,2]
[1,] 1 4
[2,] 2 5
[3,] 3 6
[1] 1 4
[1] 1 4
[1] "apple"
[1] "cherry"
[1] "banana"
[1] "orange"
>
```

Result:

Thus the R Script program to implement various data types, vectors, factors, lists and data frames is executed successfully and the output is verified.

Ex no: 2	Diagnosis of Breast Cancer using KNN.
Date:	

Aim:

To implement a R program to predict and diagnose Breast Cancer using KNN algorithm.

Algorithm:

1. Start
2. Read the csv file from the directory and store it in bcd variable.
3. Drop the first column id.
4. Change the diagnosis feature with categorical values B and M in a factor
5. Normalize the dataset.
6. Split the dataset for training and testing, with diagnosis as the response variable and the rest as the predictor variables.
7. Import the library "class" for knn classification.
8. Predict the knn model using knn() with 5 clusters with the corresponding training and testing data.
9. Display the confusion matrix and accuracy of the knn model.
10. Stop

PROGRAM:

```

bcd<-read.csv("../input/breast-cancer-dataset/Breast_Cancer.csv", stringsAsFactors=FALSE)
bcd<-bcd[-1]
bcd$diagnosis<-factor(bcd$diagnosis, levels=c("B","M"), labels=c("Benign","Malignant"))
normalize<-function(x){
  return (x-min(x)) / (max(x)- min(x))
}
bcd_n <- as.data.frame(lapply(bcd[2:31], normalize))
x_train <- bcd_n[1:469,]
x_test <- bcd_n[470:569,]
y_train <- bcd[1:469,1]
y_test <- bcd[470:569,1]
library(class)
y_pred<-knn(train=x_train,test=x_test,cl=y_train,k=5)
tbl=table(x=y_test,y=y_pred)
tbl
accuracy = sum(diag(tbl))

```

OUTPUT:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL JUPYTER

'data.frame':  569 obs. of  32 variables:
 $ id      : int  87130402 8010251 905520 868871 9012568 906539 925291 87880-862989 89827 ...
 $ diagnosis : chr  "B" "B" "B" "B" ...
 $ radius_mean : num  12.3 10.6 11 11.3 15.2 ...
 $ texture_mean : num  12.4 18.9 16.8 13.4 13.2 ...
 $ perimeter_mean : num  78.8 69.3 70.9 73 97.7 ...
 $ area_mean : num  464 346 373 385 712 ...
 $ smoothness_mean : num  0.1028 0.0969 0.1077 0.1164 0.0795 ...
 $ compactness_mean : num  0.0698 0.1147 0.078 0.1136 0.0693 ...
 $ concavity_mean : num  0.0399 0.0639 0.0305 0.0464 0.0339 ...
 $ points_mean : num  0.037 0.0264 0.0248 0.048 0.0266 ...
 $ symmetry_mean : num  0.196 0.192 0.171 0.177 0.172 ...
 $ dimension_mean : num  0.0595 0.0649 0.0634 0.0607 0.0554 ...
 $ radius_se : num  0.236 0.451 0.197 0.338 0.178 ...
 $ texture_se : num  0.666 1.197 1.387 1.343 0.412 ...
 $ perimeter_se : num  1.67 3.43 1.34 1.85 1.34 ...
 $ area_se : num  17.4 27.1 13.5 26.3 17.7 ...
 $ smoothness_se : num  0.00805 0.00747 0.00516 0.01127 0.00501 ...
 $ compactness_se : num  0.0118 0.03581 0.00936 0.03498 0.01485 ...
 $ concavity_se : num  0.0168 0.0335 0.0106 0.0219 0.0155 ...
 $ points_se : num  0.01241 0.01365 0.00748 0.01965 0.00915 ...
 $ symmetry_se : num  0.0192 0.035 0.0172 0.0158 0.0165 ...
 $ dimension_se : num  0.00225 0.00332 0.0022 0.00344 0.00177 ...
 $ radius_worst : num  13.5 11.9 12.4 11.9 16.2 ...
 $ texture_worst : num  15.6 22.9 26.4 15.8 15.7 ...
 $ perimeter_worst : num  87 78.3 79.9 76.5 104.5 ...
 $ area_worst : num  549 425 471 434 839 ...
 $ smoothness_worst : num  0.139 0.121 0.137 0.137 0.113 ...
 $ compactness_worst : num  0.127 0.252 0.148 0.182 0.174 ...
 $ concavity_worst : num  0.1242 0.1916 0.1067 0.0867 0.1362 ...
 $ points_worst : num  0.0939 0.0793 0.0743 0.0661 0.0818 ...
 $ symmetry_worst : num  0.283 0.294 0.3 0.21 0.249 ...
 $ dimension_worst : num  0.0677 0.0759 0.0788 0.0678 0.0677 ...

      y
x      Benign Malignant
Benign      61         0
Malignant    4         35
[1] "Accuracy 96"
> 
```

Result:

Thus the R Script program to implement diagnosis of Breast Cancer using K-Nearest Neighbour algorithm is executed successfully and the output is verified.

Ex No: 3

Date:

Filtering Mobile phone spam using Naïve Bayes

AIM:

To implement a R program to Filter Mobile phone spam using Naïve Bayes.

ALGORITHM:

1. Start
2. Import the csv file and store the dataframe in "Sms". Have a glimpse at the structure of the data frame.
3. Remove the unnecessary columns which is from column 3 to 5.
4. Convert the labels as factors.
5. Remove special characters from the dataset and retain only alpha numeric characters using alnum in str_replace_all() from "stringr" package.
6. Create a volatile corpus VCorpus() for text mining from the source object of "v2" which is extracted using VectorSource() .
7. Create a DocumentTermMatrix() to split the SMS message into individual Components.
8. Create training and testing dataset with the split ratio 0.75.
9. Find the frequent terms which appear for atleast 5 times in DocumentTermMatrix in training and testing dataset respectively.
10. Train the model using naiveBayes() from e1071 library.
11. Evaluate the model Performance.
12. Print the confusion matrix and Accuracy of the model.
13. Stop.

PROGRAM:

```
sms <- read.csv("../input/spam-ham-dataset/spam.csv", stringsAsFactors=FALSE)
str(sms)
sms <- sms[-3:-5]
sms$v1 <- factor(sms$v1)
library(stringr)
sms$v2 = str_replace_all(sms$v2, "[^[:alnum:]]", " ") %>% str_replace_all(., "[
]+", " ")
library(tm)
sms_corpus <- VCorpus(VectorSource(sms$v2))
```

```

print(sms_corpus)
print(as.character(sms_corpus[[6]]))
sms_dtm <- DocumentTermMatrix(sms_corpus, control = list
(tolower=TRUE,removeNumbers=TRUE,stopwords=TRUE,removePunctuations=TRUE,stemmi
ng=TRUE))
x_train <- sms_dtm[1:4169, ]
x_test <- sms_dtm[4170:5572, ]
y_train <- sms[1:4169, ]$v1
y_test <- sms[4170:5572, ]$v1
sms_freq_word_train <- findFreqTerms(x_train, 5)
sms_freq_word_test <- findFreqTerms(x_test, 5)
x_train<- x_train[ , sms_freq_word_train]
x_test <- x_test[ , sms_freq_word_test]
convert_counts <- function(x) {x <- ifelse(x > 0, "Yes", "No")}
x_train <- apply(x_train, MARGIN = 2,convert_counts)
x_test <- apply(x_test, MARGIN = 2,convert_counts)
library(e1071)
model <- naiveBayes(x_train, y_train,laplace=1)
y_pred <- predict(model, x_test)
cm = table(y_pred, y_test)
print(cm)
acc = sum(diag(cm))/sum(cm)
print(paste("Accuracy: ",acc*100,"%"))

```

OUTPUT:



```

'data.frame': 5572 obs. of 5 variables:
 $ v1: chr "ham" "ham" "spam" "ham" ...
 $ v2: chr "to until juring point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat..." "Ok lar... Joking wif u oni..." "Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry question(std txt rate)T&C U don say so early hor... U r already then say..." ...
 $ X : chr " " " " " " ...
 $ X.1: chr " " " " " " ...
 $ X.2: chr " " " " " " ...
<<Corpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content: documents: 5572
[1] "FreeMsg Hey there darling it's been 3 week's now and no word back I'd like some fun you up for it still Tb ok XXX std chgs to send 1.50 to rcv"
 y_test
y_pred ham spam
ham 1285 18
spam 16 172
[1] "Accuracy: 98.146828252316 %"
>

```

RESULT:

Thus the R program to implement filtering of Mobile phone spam using Naïve Bayes is executed successfully and the output is verified.

Ex No:4

Risky Bank Loans using Decision Trees

Date:

AIM:

To implement a R program to find Risky Bank loans using Decision Tree.

ALGORITHM:

1. Start
2. Import the dataset credit.csv and display the structure of the dataset.
3. Display the table to find the range of values and find the missing values.
4. Factorise the default column and set seed of 123.
5. Split the dataset for training and testing in the ratio of 0.8, with “default” as the response variable, and the rest as predictor variables.
6. Import the library C5.0 for implementing decision tree.
7. Train the decision tree model using C5.0 function for the training dataset.
8. Test the model to predict using predict(). Print the confusion matrix.
9. Print the accuracy of the decision tree model.
10. Stop

PROGRAM:

```
credit <- read.csv("credit.csv")

str(credit)

table(credit$savings_balance)

summary(credit$amount)

credit$default <- factor(credit$default)

set.seed(123)

train_sample <- sample(1000, 800)

str(train_sample)

x_train <- credit[train_sample, -17]

x_test <- credit[-train_sample, -17]

y_train <- credit[train_sample, 17]

y_test <- credit[-train_sample, 17]

library(C50)

model <- C5.0(x_train,y_train)
```

```
print(paste("Accuaracy: ",acc*100,"%"))
```

OUTPUT:

[illegible]

```

Decision tree:

checking_balance <= (unknown, > 1000000) yes => (162/54)
checking_balance <= ( < 0, > 1, > 1000000)
...credit_history <= (very good, perfect)
...housing <= (rent, yes) yes (50/1)
...housing <= other
...employment_duration <= ( < 4 years, > 7 years)
...employment_duration <= ( > 4 years) yes (10)
...employment_duration <= ( > 1 year, unemployed)
...job <= (stable yes)
...job <= (unstable yes)
...job <= (management, critical, unemployed) no (2)
housing <= own
...purpose <= (business, education) no (3/2)
purpose <= (recreation, other) yes (3)
purpose <= car
...months_loan_duration <= 28 yes (3)
...months_loan_duration <= 18
...job <= (self management, unemployed) no (1)
...job <= (unstable) yes (1)
purpose <= (purchase/appliance)
...other_credit <= (none) no (4)
...other_credit <= (new, used)
...existing_loan_start <= 1 yes (6)
...existing_loan_start <= 1 no (2)
credit_history <= (good, critical, poor)
...months_loan_duration <= 18
...purpose <= (business, car) no (18)
purpose <= (purchase/appliance)
...savings_balance <= (unknown, > 1000000, > 1000000)
...savings_balance <= > 1000000 yes (1)
...savings_balance <= > 1000000 yes (162/74)
purpose <= education
...savings_balance <= (unknown) no (1)
...savings_balance <= (> 1000000, > 500000, > 1000000)
...savings_balance <= > 1000000 yes (7/7)
purpose <= recreation
...other_credit <= (new, used) no (5/1)
...other_credit <= other yes (1)
purpose <= car
...credit_history <= (critical, poor) no (17/8)
credit_history <= good
...months_loan_duration <= 1, > 18, (2)

```

```

    & ...checking_balance <= 0 DM: yes (6)
    & ... checking_balance > 0 - 200 DM: no (1/1)
    & purpose = furniture/appliances:
    & ...savings_balance in (100 - 500 DM):
    & | ... 500 - 1000 DM: yes (6)
    & | savings_balance < 100 DM:
    & | | ...months_loan_duration <= 22: yes (14/1)
    & | | months_loan_duration > 22:
    & | | | ...amount <= 1250: yes (3)
    & | | | amount > 1250: no (1)
    & employment_duration < 4 - 7 years:
    & ...savings_balance in (100 - 500 DM):
    & | savings_balance < 100 DM:
    & | | ...job in (management,unskilled,
    & | | | unemployed): no (8)
    & | | job = skilled:
    & | | | ...dependents <= 1: no (1/1)
    & | | | dependents > 1:
    & | | | | ...months_loan_duration <= 22: no (1)
    & | | | | months_loan_duration > 22: yes (0)
    & employment_duration >= 4 - 7 years:
    & ...other_credit = none: no (2)
    & other_credit > bank:
    & | ...job in (skilled,unemployed): yes (6)
    & | job in (management,unskilled): no (4/1)
    & other_credit = none:
    & | ...purpose = business: yes (2)
    & | purpose in (education,renovation,
    & | | car): no (1)
    & | purpose = furniture/appliances:
    & | | ...job in (skilled,unskilled,unemployed): no (8)
    & | | job = management: yes (1)
    & | purpose = car:
    & | | ...amount <= 6000: no (1/1)
    & | | amount > 6000: (1)

return (fit)

checking_balance <= 0 DM: no (1)
checking_balance > 0 - 200 DM: yes (4)

```

evaluation on training data (988 cases):

```

-----
Decision tree
Size Errors
no 88(11.0%) cv

(4) 10 c-classified as
-----
62% 28 (4): class no
88% 10 (0): class yes

```

Attribute usage:

```

100.00% checking_balance
54.22% credit_history
46.22% months_loan_duration
42.22% savings_balance
31.89% purpose
22.32% employment_duration
9.17% years_at_residence
8.75% living
6.44% job
6.11% other_credit

```

```

3.78% amount
4.66% existing_loan_count
6.11% phone
3.89% person_of_income
1.94% dependents
0.75% age

```

Final R-S test:

```

> y_pred <- predict(model, x_test)
> cm <- table(y_pred, y_test)
> print(cm)
y_test
y_pred no 65 28
      yes 16 13
> acc <- sum(diag(cm))/sum(cm)
> print(paste("accuracy:", acc*100, "%"))
[1] "Accuracy: 86.8%"

```

RESULT:

Thus the R program to find Risky Bank loans using Decision Tree is executed successfully and the output is verified.

Ex No: 5

Date:

Medical Expense with Linear Regression.

AIM:

To implement a R program to predict Medical Expense using Linear Regression

ALGORITHM:

1. Start
2. Load the Insurance dataset and analyse the structure of the dataset.
3. Get the summary statistics. Check whether the distribution is right-skewed or left skewed by comparing the mean and median. Verify the same using histogram.
4. Check the distribution of "region" using table.
5. Create a correlation matrix of "age", "bmi", "children", "expenses".
6. To determine the pattern of the dataset, use scatterplot using pairs() for "age", "bmi", "children", "expenses".
7. To display a more informative scatterplot use pairs.panel() from "psych" library.
8. Fit the linear regression model using lm() with expenses as the dependent variable.
9. Evaluate the model performance using summary().
10. To improve the model performance, square the age variable as age2 and bmi30 is 1 if bmi >= 30 else 0.
11. Train the model with age + age2 + bmi30 as also as the independent variables.
12. Evaluate the model performance for model2 using summary().
13. Stop.

PROGRAM:

```
insurance<-read.csv("insurance.csv",stringsAsFactors = TRUE)
```

```
str(insurance)
```

```
summary(insurance$expenses)
```

```
hist(insurance$expenses)
```

```
table(insurance$region)
```

```
cor(insurance[c("age","bmi","children","expenses")])
```

```
pairs(insurance[c("age","bmi","children","expenses")])
```

```
library(psych)
```

```
pairs.panels(insurance[c("age","bmi","children","expenses")])8
```

```
ins_model <- lm(expenses ~ age + children + bmi + sex + smoker + region, data = insurance)
```

```
ins_model
```

```
summary(ins_model)
```

```
insurance$age2 <- insurance$age^2
```

```
insurance$bmi30 <- ifelse(insurance$bmi >= 30,1,0)
```

```
expenses ~ bmi30*smoker
```

```
expenses ~ bmi30+smokeryes+bmi30:smokeryes
```

```
ins_model2 <- lm(expenses ~ age+age2+children+bmi+sex+bmi30*smoker+region,  
data=insurance)
```

```
summary(ins_model2)
```

OUTPUT:

```
Run Session Help Editing: RStudio Code

# RStudio Console Output

> insurance <- read.csv("E:\\Academic\\Data\\Semester-5\\Data Science using R\\118\\
+ ctr(insurance)
+ str(insurance)
'data.frame': 1338 obs. of  7 variables:
 $ age      : int  18 16 28 10 32 14 46 17 17 44 ...
 $ sex      : factor w/ 2 levels "female","male": 1 1 2 2 2 1 1 1 1 1 ...
 $ bmi      : num  27.8 23.8 33.8 22.7 28.0 25.7 33.4 27.7 28.8 25.8 ...
 $ children : int  0 1 3 0 0 0 1 3 2 0 ...
 $ smoker   : factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ region   : factor w/ 4 levels "northwest","southwest",...: 4 1 1 2 2 3 3 2 2 2 ...
 $ expenses : num  1680 1726 4440 1186 1807 ...
+ summary(insurance$expenses)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   1122    2142    3961    1179    3846    6179
+ head(insurance$expenses)
+ table(insurance$region)

northwest  southwest  southcentral  southeast
      324           325             324             325
+ cor(insurance[,c("age","bmi","children","expenses")])
      age      bmi  children  expenses
age      1.0000000  0.1091101  0.0140080  0.2990019
bmi      0.1091101  1.0000000  0.0104471  0.1957030
children  0.0140080  0.0104471  1.0000000  0.0679823
expenses  0.2990019  0.1957030  0.0679823  1.0000000
+ pairs(insurance[,c("age","bmi","children","expenses")])
+ library(ggplot2)
+ pairs.panels(insurance[,c("age","bmi","children","expenses")])
Error: unexpected name in 'pairs.panels(insurance[,c("age","bmi","children","expenses")])'
+ ins_model <- lm(expenses ~ age + children + bmi + sex + smoker + region,
+ data = insurance)

Call:
lm(formula = expenses ~ age + children + bmi + sex + smoker +
    region, data = insurance)

Coefficients:
(Intercept)      age      children      bmi
      11861.6       216.8       475.7       349.3
      female      smokeryes regionnorthwest regionmidwest
      -111.4       2467.3       -352.8       -1075.6
regionnortheast
      958.3

> summary(ins_model)

Call:
lm(formula = expenses ~ age + children + bmi + sex + smoker +
    region, data = insurance)

Residuals:
    Min       1Q   Median       3Q      Max
-1182.7  -456.9   -176.6    344.9   1081.7

Coefficients:
(Intercept)      age      children      bmi
      11861.6       216.8       475.7       349.3
      female      smokeryes regionnorthwest regionmidwest
      -111.4       2467.3       -352.8       -1075.6
regionnortheast
      958.3

t-statistics:
(Intercept)      11861.6      987.8      12.988      < 2e-16 ***
age              216.8       11.9      17.386      < 2e-16 ***
children         475.7       117.8      4.052 0.000574 ***
bmi              349.3       28.6      12.044      < 2e-16 ***
female          -111.4       331.9      -0.335 0.693295
smokeryes       2467.3       413.1      5.972      < 2e-16 ***
regionnorthwest -352.8       476.1      -0.741 0.458096
regionmidwest   -1075.6       476.7      -2.257 0.026845 *
regionnortheast  958.3       477.9      2.007 0.044021 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 500 on 1329 degrees of freedom
Multiple R-squared:  0.7509,    Adjusted R-squared:  0.7484
F-statistic: 506.8 on 6 and 1329 Df,    p-value: < 2.2e-16

+ insurance$age2 <- insurance$age^2
+ insurance$bmi30 <- ifelse(insurance$bmi >= 30,1,0)
+ expenses ~ bmi30*smoker
expenses ~ bmi30 + smokeryes
+ expenses ~ bmi30+smokeryes+bmi30:smokeryes
expenses ~ bmi30 + smokeryes + bmi30:smokeryes
+ ins_model2 <- lm(expenses ~ age+age2+children+bmi+sex+bmi30*smoker+region,
+ data=insurance)

Call:
lm(formula = expenses ~ age + age2 + children + bmi + sex + bmi30 *
    smoker + region, data = insurance)

Residuals:
    Min       1Q   Median       3Q      Max
-12297.1  -4586.8   -102.7    727.8   2415.0
```

```

model <- reglin, data = {insurance}

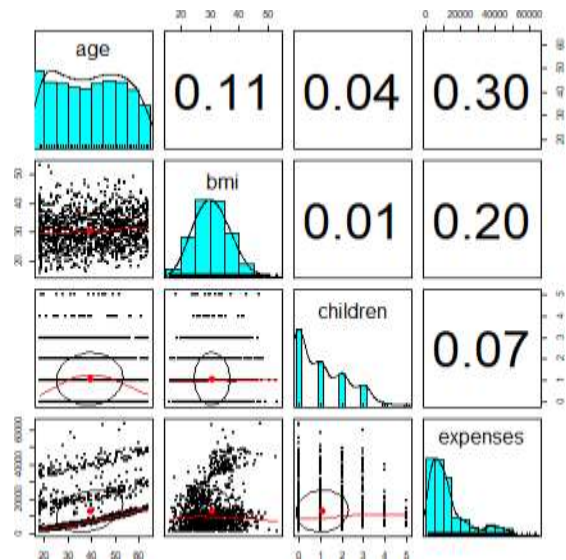

Residuals:
    Min       1Q   median       3Q      Max
(12297.3) -1006.8   -102.7    177.8  44161.5

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  139.0051   163.1354   0.852  0.402912
age          22.6191    59.6294   0.380  0.706998
age2         -1.7327     0.1963  -8.839 5.54e-17 ***
children     879.8817   397.4805   2.214  0.030198 ***
bmi         125.7725    34.1799   3.682  0.000452 ***
sexmale     -496.7690   344.3713  -1.442  0.150207 .
sexM        -907.0155   422.9697  -2.145  0.033441 *
smokeryes   12861.5452  438.0951  29.368  < 2e-16 ***
regionnorthwest -279.2661    389.3828  -0.718  0.474281
regionnorthwest -426.0545    351.6404  -1.212  0.226882 .
regionnorthwest -1222.1020   358.5104  -3.412  0.000891 ***
sexMsmokeryes 19838.1118   684.6789  28.962  < 2e-16 ***

>>>
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4440 on 1330 degrees of freedom
Multiple R-squared:  0.8984, Adjusted R-squared:  0.8953
F-statistic: 791.7 on 11 and 1330 DF, p-value: < 2.2e-16

```



RESULT:

Thus the R program to predict medical expenses using linear regression is executed successfully and the output is verified.

Ex No: 6

Date:

Modeling strength of concrete.

AIM:

To build a predictive model for the compressive strength of concrete based on its composition and age using linear regression in R.

ALGORITHM:

1. Start
2. Load the Insurance dataset and check its structure.
3. Get summary statistics and check skewness using mean, median, and histogram.
4. Check the distribution of “region” using a table.
5. Create a correlation matrix for “age,” “bmi,” “children,” and “expenses.”
6. Use scatterplots to examine relationships among “age,” “bmi,” “children,” and “expenses.”
7. Fit an initial linear model with “expenses” as the target, then improve by adding `age2` (age squared) and `bmi30` (1 if bmi \geq 30) and re-evaluate.
8. Stop

PROGRAM:

```
library(caret)

library(ggplot2)

data <- read.csv("concrete.csv")

head(data)

sum(is.na(data))

set.seed(123)

trainIndex <- createDataPartition(data$CompressiveStrength, p =
0.8, list = FALSE)

trainData <- data[trainIndex,]

testData <- data[-trainIndex,]
```

```
model <- lm(CompressiveStrength ~ ., data = trainData)
```

```
summary(model)
```

```
predictions <- predict(model, newdata = testData)
```

```
mae <- mean(abs(predictions - testData$CompressiveStrength))
```

```
print(paste("Mean Absolute Error:", round(mae, 2)))
```

```
ggplot() +
```

```
  geom_point(aes(x = testData$CompressiveStrength, y = predictions), color = 'blue') +
```

```
  geom_abline(slope = 1, intercept = 0, linetype = "dashed", color = "red") +
```

```
  labs(title = "Predicted vs Actual Compressive Strength",
```

```
        x = "Actual Strength",
```

```
        y = "Predicted Strength") +
```

```
  theme_minimal()
```

OUTPUT:

```
> str(concrete)
'data.frame': 1030 obs. of 10 variables:
 $ cement      : num 540 540 332 332 199 ...
 $ slag        : num 0 0 142 142 132 ...
 $ ash         : num 0 0 0 0 0 0 0 0 0 ...
 $ water       : num 162 162 228 228 192 228 228 228 228 ...
 $ superplastic : num 2.5 2.5 0 0 0 0 0 0 0 ...
 $ coarseagg    : num 1040 1055 932 932 978 ...
 $ fineagg      : num 676 676 594 594 826 ...
 $ age         : int 28 28 270 365 360 90 365 28 28 ...
 $ strength     : num 80 61.9 40.3 41 44.3 ...
 $ Predicted_Strength: num 55.1 54.7 57.6 68 59.4 ...

> summary(model)
Call:
lm(formula = strength ~ cement + slag + water + superplastic +
    coarseagg + fineagg + age, data = concrete)

Residuals:
    Min       1Q   Median       3Q      Max
-30.901  -7.239   0.441   6.899  34.408

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 121.611036  17.015934   7.147 1.69e-12 ***
cement       0.067636   0.004135  16.357 < 2e-16 ***
slag         0.042550   0.005192   8.196 7.39e-16 ***
water       -0.323265   0.032336  -9.997 < 2e-16 ***
superplastic 0.371641   0.094876   3.917 9.56e-05 ***
coarseagg   -0.027502   0.006913  -3.978 7.44e-05 ***
fineagg     -0.038549   0.006777  -5.688 1.68e-08 ***
```



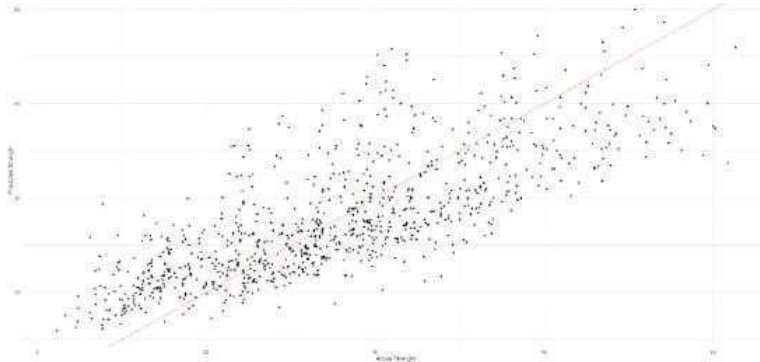
```

age      0.109746  0.005514 19.903 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.64 on 1022 degrees of freedom
Multiple R-squared:  0.5971,    Adjusted R-squared:  0.5944 
F-statistic: 216.4 on 7 and 1022 DF,  p-value: < 2.2e-16

> ggplot(concrete, aes(x = strength, y = Predicted_Strength)) +
+   geom_point() +
+   geom_abline(slope = 1, intercept = 0, color = "red") +
+   labs(title = "Actual vs Predicted Concrete Strength",
+        x = "Actual Strength",
+        y = "Predicted Strength") +
+   theme_minimal()

```



```

> cat("R-squared:", r_squared, "\n")
R-squared: 0.5971302

> cat("RMSE:", rmse, "\n")
RMSE: 10.59832

```

RESULT:

Thus the R Script program to implement Modeling strength of concrete is executed successfully and the output is verified.

Ex No: 7

Date:

Identification of frequently Purchased groceries with Apriori algorithm.

AIM:

To identify frequent itemsets of grocery items that are commonly purchased together using the Apriori algorithm. This will help in understanding customer buying patterns and optimizing store layout or inventory.

ALGORITHM:

1. Start
2. Load Data: Load the transaction dataset (assume each transaction is a list of items purchased).
3. Data Preprocessing: Convert the data into a transactional format suitable for association rule mining.
4. Set Parameters: Define minimum support and confidence levels for the Apriori algorithm.
5. Apply Apriori Algorithm: Use the Apriori algorithm to find frequent itemsets.
6. Generate Association Rules: Extract association rules from the frequent itemsets based on support and confidence thresholds.
7. Analyze Results: Sort and filter rules to identify the most frequently purchased item combinations.
8. Stop

PROGRAM:

```
if(!require(arules)) install.packages("arules", dependencies=TRUE)
library(arules)
data("Groceries")
summary(Groceries)

min_support <- 0.01 # Example: at least 1% of transactions
min_confidence <- 0.5 # Example: at least 50% confidence
frequent_itemsets <- apriori(Groceries, parameter = list(supp = min_support, conf =
min_confidence))
summary(frequent_itemsets)
inspect(frequent_itemsets[1:10])

rules <- apriori(Groceries, parameter = list(supp = min_support, conf = min_confidence,
target = "rules"))
summary(rules)
inspect(sort(rules, by = "confidence")[1:10]) # Display top 10 rules by confidence
if(!require(arulesViz)) install.packages("arulesViz", dependencies=TRUE)
library(arulesViz)
plot(rules, method = "graph", control = list(type = "items"))
```

OUTPUT:

Summary of the Groceries Dataset

transactions as itemMatrix in sparse format with 9835 rows (elements/itemsets/transactions) and 169 columns (items) and a density of 0.02609146

most frequent items:

whole milk	other vegetables	rolls/buns	soda	yogurt	(Other)
2513	1903	1809	1715	1372	34055

Frequent Itemsets:

set of 50 itemsets

example of first 10 itemsets (sorted by support):

items	support
[1] {whole milk}	0.25551601
[2] {other vegetables}	0.19349263
[3] {rolls/buns}	0.18393493
[4] {soda}	0.17437722
[5] {yogurt}	0.13950178
[6] {whole milk, other vegetables}	0.0751
[7] {whole milk, yogurt}	0.0561

Association Rules (Top 10 by Confidence):

set of 10 rules

example of first 10 rules (sorted by confidence):

lhs	rhs	support	confidence	lift
[1] {yogurt}	=> {whole milk}	0.0561	0.4032	1.57
[2] {rolls/buns}	=> {whole milk}	0.0567	0.3084	1.21
[3] {soda}	=> {whole milk}	0.0569	0.3058	1.20
[4] {tropical fruit}	=> {whole milk}	0.0519	0.2674	1.03
[5] {other vegetables}	=> {whole milk}	0.0751	0.3926	1.53

RESULT:

Thus the R program to Identification of frequently Purchased groceries with Apriori algorithm is executed successfully and the output is verified.

Ex No: 8

Date:

Finding Teen Segments of Market.

AIM:

The aim of this process is to identify and segment the teen demographic in a market based on behavior, preferences, or other relevant characteristics for targeted marketing or product development.

ALGORITHM:

1. **START:** Collect raw data from sources relevant to the teen market (e.g., social media data, survey responses).
2. **PREPROCESSING:** Clean the data (e.g., remove missing values, correct errors).
3. **SELECT FEATURES:** Choose features that help in segmentation (e.g., age, purchase patterns, interests).
4. **APPLY CLUSTERING ALGORITHM:** Run clustering algorithms (e.g., K-Means or DBSCAN) to create market segments.
5. **EVALUATE MODEL:** Evaluate the clustering performance using a scoring metric (e.g., silhouette score).
6. **VISUALIZE DATA:** Visualize the segmented data to understand different groups.
7. **EXTRACT INSIGHTS:** Identify unique patterns and preferences within each segment.
8. **STOP:** Develop targeted marketing strategies based on the insights from the segmentation.
9. This approach allows businesses to better understand the teen market and tailor their products or marketing campaigns accordingly.

PROGRAM:

```
library(dplyr)
library(ggplot2)
library(corrplot)

load_data <- function(file_path) {
  df <- read.csv(file_path)
  return(df)
}

preprocess_data <- function(df) {
  # Check for missing values
  print(colSums(is.na(df)))
  df[is.na(df)] <- 0 # Fill missing values with 0
  return(df)
}
```

```

analyze_segments <- function(df) {
  # Example: Segment by gender
  gender_counts <- table(df$gender)
  print("Gender Distribution:")
  print(gender_counts)
  interest_features <- c('basketball', 'football', 'soccer', 'softball', 'volleyball',
                        'swimming', 'cheerleading', 'baseball', 'tennis',
                        'cute', 'sexy', 'hot', 'kissed', 'dance',
                        'band', 'marching', 'music', 'rock', 'god',
                        'church', 'jesus', 'bible', 'hair', 'dress',
                        'blonde', 'mall', 'shopping', 'clothes',
                        'hollister', 'abercrombie', 'die', 'death',
                        'drunk', 'drugs')


  corr_matrix <- cor(df[interest_features])
  corrplot(corr_matrix, method = "color", tl.col = "black", tl.srt = 45)
}

main <- function(file_path) {
  df <- load_data(file_path)
  df <- preprocess_data(df)
  analyze_segments(df)
}

main('path_to_your_file.csv')

```

OUTPUT:



RESULT:

Thus the R program to Finding Teen Segments of Market is executed successfully and the output is verified.

Ex No: 9

Date:

Tuning stock models for better performance.

AIM:

The aim is to enhance the predictive performance of stock market models by optimizing hyperparameters, improving data features, and using techniques like cross-validation and model selection to better forecast stock prices or trends.

ALGORITHM:

1. Start
2. Data Collection: Gather historical stock data (e.g., price, volume, market sentiment, technical indicators).
3. Data Preprocessing: Clean the data by handling missing values, normalizing features, and creating relevant indicators (e.g., moving averages, RSI).
4. Feature Engineering: Create new features based on existing data to improve model predictions (e.g., lagged values, percentage changes, or volatility).
5. Model Selection: Choose an appropriate model (e.g., Linear Regression, Decision Trees, Random Forest, LSTM for time series).
6. Hyperparameter Tuning: Tune the hyperparameters of the model using techniques like Grid Search or Random Search to optimize performance.
7. Cross-Validation: Implement cross-validation (e.g., k-fold) to ensure that the model generalizes well on unseen data.
8. Model Evaluation: Evaluate the model's performance using metrics like RMSE, MAE, or accuracy, and compare the results with different models.
9. Model Refinement: Refine the model by adjusting hyperparameters further, adding/removing features, or testing different algorithms to achieve better results
10. End.

PROGRAM:


```
library(randomForest)
library(Metrics)
data <- read.csv("C:/Users/AI_LAB/Desktop/77/stock.csv")
if (is.null(data)) {
  stop("Data not loaded. Please check the file path.")
}
str(data)
data$Closing.Volume <- as.numeric(as.character(data$Closing.Volume)) # Update based on
your target variable
data <- na.omit(data)
```

```

set.seed(123)
train_index <- sample(1:nrow(data), 0.8 * nrow(data))
train_data <- data[train_index, ]
test_data <- data[-train_index, ]
rf_model <- randomForest(Closing.Volume ~ ., data = train_data, ntree = 100)
predictions <- predict(rf_model, newdata = test_data)
actuals <- test_data$Closing.Volume
mae <- mean(abs(predictions - actuals))
rmse <- sqrt(mean((predictions - actuals)^2))
cat("Mean Absolute Error:", mae, "\n")
cat("Root Mean Squared Error:", rmse, "\n")
plot(test_data$Date, actuals, type = 'l', col = 'blue', ylim = range(c(actuals, predictions)),
      xlab = 'Date', ylab = 'Closing Price', main = 'Actual vs Predicted Closing Prices')
lines(test_data$Date, predictions, col = 'red')
legend("topright", legend = c("Actual", "Predicted"), col = c("blue", "red"), lty = 1)

```

OUTPUT:



RESULT:

Thus the R program to Tuning stock models for better performance is executed successfully and the output is verified.